

**CALIBRATION-AWARE RECURRENT NEURAL
NETWORKS
FOR MULTIVARIATE PROBABILISTIC
FORECASTING
UNDER DISTRIBUTION SHIFT**

by
Jonathan Ma

A thesis submitted to The Johns Hopkins University in conformity
with the requirements for the degree of M.S.E. Applied Mathematics and Statistics

Baltimore, Maryland
Nov 2026

© 2026 Jonathan Ma
All rights reserved

Abstract

Probabilistic forecasts are useful only when their stated uncertainty is credible. Likelihood-trained recurrent neural networks may be miscalibrated, while Bayesian model averaging is not guaranteed to remain calibrated when the likelihood, prior, or conditional dynamics are misspecified. This paper develops the Calibration-Aware Recurrent Neural Network (CA-RNN), adapting the calibration-aware learning framework of Huang, Park, and Simeone and existing differentiable regression and multivariate pre-rank regularization to sequential multivariate forecasting. The training criterion augments joint negative log likelihood with a differentiable discrepancy between projected probability integral transforms (PITs) and the uniform distribution and a temporal penalty on selected lagged PIT moments. CA-RNN is a conditional stochastic forecasting model: its recurrent state parameterizes a predictive distribution from which future observations can be sampled. The primary empirical comparison is between RNN and CA-RNN, with Bayesian variants retained as a supporting replication and VAR–GARCH as the principal econometric benchmark. The evaluation then extends to model misspecification and distribution shift. In a post-protocol simulation with 50 independent datasets and an isolated innovation-scale shift, CA-RNN and RNN perform similarly at zero shift, whereas CA-RNN has lower Energy Score, variogram score, interval score, test negative log likelihood, and analytic projected-PIT calibration error at both shifted severities. A frozen replication on a disjoint four-asset panel confirms a smaller projected-PIT calibration improvement, but finds significantly worse MSE, Energy Score, interval score, and PIT temporal dependence than the likelihood-trained RNN. A VAR with marginal GARCH(1, 1) scales and joint

Abstract

standardized-residual resampling provides substantially better interval calibration than either neural model. The results show that directly optimizing a calibration diagnostic can generalize to new assets without improving the predictive distribution as judged by proper scores. Supporting results prove that the finite calibration criterion is non-identifying, construct dependent uniform PITs with zero covariance, and express Energy and interval-score regret in quantities that the calibration loss does not control.

Keywords: calibration, recurrent neural networks, multivariate forecasting, probability integral transform, distribution shift, financial time series

Primary reader and thesis advisor

Dr. Luhao Zhang
Assistant Professor
Department of Applied Mathematics and Statistics
Johns Hopkins University, Baltimore, MD

Table of Contents

Abstract	ii
List of Tables	vii
List of Figures	ix
Chapter 1 Introduction	1
1.1 Motivation	1
1.2 Research questions	2
1.3 Contributions and novelty boundary	3
1.4 Scope and non-claims	3
Chapter 2 Related Work	4
2.1 Calibration-aware Bayesian learning	4
2.2 Calibration, sharpness, and proper scoring	4
2.3 Density-forecast diagnostics under temporal dependence	5
2.4 Calibration-aware neural prediction	6
2.5 Multivariate pre-rank calibration and regularization	7
2.6 Probabilistic recurrent forecasting, misspecification, and shift	8
2.7 Positioning of the present study	9
Chapter 3 Forecasting Setup and CA-RNN Methodology	10
3.1 Forecasting setup and calibration	10
3.1.1 Sequential prediction problem	10
3.1.2 Joint Gaussian predictive head	10
3.1.3 Negative log likelihood	11
3.1.4 Projected probability integral transforms	12
3.2 Calibration-aware RNN methodology	14

Table of Contents

3.2.1	Recurrent architecture	14
3.2.2	Differentiable projected calibration loss	15
3.2.3	Temporal PIT-moment loss	16
3.2.4	Frequentist CA-RNN objective	18
3.2.5	Supporting Bayesian extension	18
3.2.6	Why targeted calibration need not improve proper scores . . .	20
3.2.7	Training algorithms	26
Chapter 4	Distribution Shift and Experimental Design	28
4.1	Model misspecification and distribution shift	28
4.1.1	Distinguishing misspecification from shift	28
4.1.2	Stress families	29
4.1.3	Robustness estimands	30
4.2	Experimental design	31
4.2.1	Primary comparison and supporting Bayesian replication . . .	31
4.2.2	Conventional benchmark ladder	31
4.2.3	Frozen external-panel replication	32
4.2.4	Calibration-weight experiment	33
4.2.5	Temporal-penalty ablation	33
4.2.6	Controlled simulations	34
4.2.7	Financial application	34
4.2.8	Status of the evidence	35
4.2.9	Paired uncertainty design	36
4.2.10	Metrics and figures	37
Chapter 5	Results and Conclusions	38
5.1	Results	38
5.1.1	Implementation validation	38
5.1.2	Primary comparison and supporting Bayesian results	38
5.1.3	Calibration-weight sensitivity	39
5.1.4	Temporal-penalty ablation	39

Table of Contents

5.1.5	Misspecification and distribution shift	40
5.1.6	Financial results	41
5.1.7	Post-protocol financial evaluation audit	42
5.1.8	Conventional-baseline results	43
5.1.9	Frozen external-panel replication	45
5.1.10	Rolling market-regime results	46
5.1.11	Interval-score mechanism under temporal shift	47
5.2	Discussion	49
5.2.1	Interpretation	49
5.2.2	Limitations	51
5.2.3	Future work	52
5.3	Conclusion	53
	Bibliographic references	55
	Appendix A Tables and Figures	58
A.1	Tables	58
A.2	Figures	62
	Appendix B Reproducibility and Post-Protocol Evaluation Audit	64
B.1	Data, transformations, and chronological partitions	64
B.2	Model, randomization, and selection settings	65
B.3	Evaluation definitions and protocol status	66

List of Tables

Table A.1	Positioning relative to the closest calibration literature. “Shift” means an explicit out-of-distribution or regime evaluation, not ordinary held-out testing.	58
Table A.2	Mapping from calibration-aware classification to forecasting. . .	58
Table A.3	Loss-component ablation.	59
Table A.4	Isolated Gaussian innovation-scale shift over 50 independent paired realizations. Each entry is the mean CA–RNN-minus-RNN difference with its realization-level Monte Carlo standard error in parentheses; negative values favor CA–RNN.	59
Table A.5	Mean financial performance over ten matched optimization seeds. Lower is better except for coverage, whose nominal target is 0.90.	59
Table A.6	Original-panel post-protocol evaluation audit: means across ten matched neural seeds. All entries are losses or diagnostic discrepancies for which lower is better. MSE and RMSE are distinct summaries.	59
Table A.7	Original-panel neural and VAR–GARCH results. Neural entries are ten-seed means; the traditional model is fitted once. Coverage targets 0.90.	60
Table A.8	Training-sample diagnostics for the selected VAR(1)–GARCH benchmark. $p_{z,20}$ and $p_{z^2,20}$ are descriptive Ljung–Box p -values at lag 20; kurtosis is the ordinary, rather than excess, standardized-residual kurtosis.	60
Table A.9	Frozen MSFT–BAC–CVX–COST replication. Neural entries are ten-seed means; coverage targets 0.90.	60
Table A.10	Five-seed regime means for RNN (R) and CA-RNN (C).	61

Table A.11	Decomposition of the interval-score difference between CA-RNN and RNN. Negative total differences favor CA-RNN.	61
-------------------	---	----

List of Figures

Figure A.1	Paired CA–RNN-minus-RNN differences across scale severity. Error bars show ± 1.96 Monte Carlo standard errors across independent DGP realizations; they do not cover uncertainty over alternative DGP families or tuning choices.	62
Figure A.2	Mean Energy Score, interval score, and projected-PIT calibration error across five matched seeds and four chronological test regimes. Lower is better.	63
Figure A.3	Exact decomposition of CA-RNN-minus-RNN interval-score differences. Tail reductions offset widening in the pandemic and inflation/tightening periods, but not in the recent period. . . .	63

Chapter 1

Introduction

1.1 Motivation

A point forecast reports one possible future; a probabilistic forecast reports a distribution of possible futures. The latter is operationally meaningful only if events assigned probability α occur at approximately rate α , while the distribution remains sufficiently sharp to discriminate among outcomes. Modern neural forecasters can optimize likelihood successfully without achieving reliable finite-sample calibration. Bayesian learning represents parameter uncertainty, but its calibration guarantees depend on idealized assumptions about the likelihood, prior, posterior computation, and data-generating mechanism.

Huang, Park, and Simeone propose calibration-aware Bayesian neural networks (CABNNs), combining a data-dependent calibration regularizer with the data-independent KL regularizer of variational Bayesian learning [1]. Their experiments compare frequentist, Bayesian, calibration-aware frequentist, and calibration-aware Bayesian classifiers; vary the calibration weight; and report expected calibration error and reliability diagrams. Their conclusion identifies evaluation under distribution shift as a direction for future work. This paper transfers that methodological template to continuous multivariate sequential prediction and makes misspecification and distribution shift a central part of the evaluation.

1.2 Research questions

The primary research question is:

How does augmenting a likelihood-trained Gaussian RNN with differentiable projected-PIT calibration penalties affect the calibration and predictive quality of one-step-ahead multivariate financial forecasts, both across chronological market regimes and relative to likelihood-only neural and VAR-GARCH benchmarks?

This question is decomposed into four focused empirical questions:

- RQ1. Across the prespecified calibration-weight grid, what tradeoff arises between projected calibration error and strictly proper predictive scores, marginal interval score, and point-forecast error?
- RQ2. What does removing CA-RNN's temporal PIT-moment penalty reveal about its effect on held-out sequential calibration error, Energy Score, marginal interval score, coverage, and point-forecast error?
- RQ3. Does CA-RNN add forecast value beyond the likelihood-only RNN and VAR-GARCH on the development and frozen external asset panels, and how does its performance relative to the RNN vary across four chronological market regimes?
- RQ4. In a replicated experiment that changes only test-innovation scale, how do the paired CA-RNN-minus-RNN differences in analytic projected-PIT error, test NLL, Energy Score, variogram score, interval score, and MSE vary with shift

severity?

1.3 Contributions and novelty boundary

This study adapts differentiable projected-PIT regularization to a recurrent conditional-density model trained on ordered multivariate observations. Its sequential extension adds lagged PIT-moment restrictions. The empirical contribution is a frozen comparison with likelihood-only recurrent and VAR–GARCH forecasts, including an isolated scale-shift simulation, chronological regimes, external-panel replication, and proper-score checks. Neither differentiable PIT regularization for regression nor multivariate pre-rank regularization is introduced here as a general new method; both have direct precedents [2–4].

1.4 Scope and non-claims

Finite projected calibration is weaker than complete multivariate distributional calibration. Calibration of a finite test set is also not a guarantee of conditional calibration at every history. We therefore avoid claims that a small collection of projections proves joint correctness. The proposed regularizer is an explicit, diagnostically aligned training objective whose value must be established empirically with proper scores and out-of-sample tests.

Chapter 2

Related Work

2.1 Calibration-aware Bayesian learning

The organizing reference is [1]. It distinguishes data-dependent calibration regularization from the data-independent prior regularization used by variational Bayesian inference. Its CA-FNN objective adds a differentiable approximation of expected calibration error to cross-entropy, and its CA-BNN objective places the same calibration penalty inside the variational expectation while retaining a KL penalty. The paper uses Gaussian variational parameters and the reparameterization trick, and its main comparison comprises FNN, BNN, CA-FNN, and CA-BNN. The reported experiments emphasize calibration-weight sweeps and reliability diagrams.

2.2 Calibration, sharpness, and proper scoring

Probabilistic calibration concerns the statistical consistency between forecast distributions and realized outcomes. Gneiting, Balabdaoui, and Raftery [5] organize forecast evaluation around the principle that a predictive distribution should be as sharp as possible subject to calibration. This distinction is central here: a diffuse predictive law can achieve acceptable empirical coverage without being informative, while a sharp but systematically incorrect law is not credible. Calibration diagnostics must therefore be accompanied by measures that reward the entire predictive distribution.

Proper scoring rules provide that complement. A scoring rule is proper when its expected value is optimized by reporting the data-generating distribution and is strictly proper when that optimum is unique [6]. Negative log predictive density, the energy score, and the interval score consequently answer a different question from a finite calibration penalty: they assess the quality of the reported distribution rather than agreement with selected moment conditions. For multivariate quantities, the energy score extends the continuous ranked probability score and supports sample-based evaluation [7]. It can, however, have limited sensitivity to misspecified dependence; variogram scores were developed in part to expose such errors [8]. The manuscript therefore reports proper scores alongside projected calibration, marginal intervals, and portfolio-tail risk rather than interpreting any one diagnostic as sufficient.

More recent work formalizes conditional calibration and links forecast evaluation to regression diagnostics [9]. This is an important qualification for the present method. CA-RNN enforces a finite set of unconditional empirical restrictions on projected PITs; it does not establish auto-calibration conditional on every forecast history. The distinction motivates both the non-identification results in Section 3.2.6 and the use of genuinely out-of-sample proper-score comparisons.

2.3 Density-forecast diagnostics under temporal dependence

Probability integral transforms (PITs) convert a correctly specified continuous univariate predictive distribution into uniform observations. Diebold, Gunther, and Tay [10] made PIT histograms and dependence diagnostics a practical framework for evaluating financial density forecasts. Berkowitz [11] proposed likelihood-based tests

after a normal transformation of the PIT, allowing uniformity and independence to be examined jointly. For overlapping or multi-step forecasts, however, transformed observations may be serially correlated even under the null. Knüppel [12] develops raw-moment tests that explicitly accommodate this dependence.

These results explain the design of the CA–RNN penalty. Marginal projected PIT moments target distributional shape, while lagged PIT-product moments target one limited form of sequential dependence. The latter is diagnostically relevant, but neither zero sample autocovariance nor uniformity on a finite grid proves independence. Consequently, sequential calibration error is treated as a targeted diagnostic, not as a complete specification test.

2.4 Calibration-aware neural prediction

Modern neural networks can produce poorly calibrated probabilities even when their predictive accuracy is high. Guo et al. [13] document this phenomenon in classification and show the effectiveness of post-hoc temperature scaling. Kumar, Sarawagi, and Jain [14] instead construct a differentiable kernel-based calibration measure that can be included during training. For continuous outcomes, Kuleshov, Fenner, and Ermon [15] propose a post-hoc recalibration procedure for predictive cumulative distribution functions and demonstrate it with deep regression and recurrent models. These papers establish that calibration can be measured or corrected in neural systems. In-training regression calibration also has direct precedents: quantile regularization penalizes departures from calibrated predictive quantiles [2]; Dheur and Ben Taieb study differentiable regression regularizers, including a smoothed empirical PIT-CDF

discrepancy [3], and subsequently integrate recalibration into neural regression training [16]. A differentiable PIT-based loss is therefore not, by itself, a novel contribution of CA-RNN. Related quantile-regression work explicitly optimizes a calibration–sharpness tradeoff and provides another training-time continuous-outcome precedent [17].

Huang, Park, and Simeone [1] bring calibration directly into Bayesian neural learning. Their calibration-aware objective combines a data-dependent differentiable calibration loss with a data-independent variational KL regularizer. CA-RNN preserves this organizing idea but changes the statistical object: class probabilities are replaced by multivariate one-step-ahead predictive distributions, and calibration is evaluated through projected PIT and sequential moment restrictions. This adaptation is not mechanical. Forecast observations are ordered and dependent, predictive covariance matters, and a multivariate law cannot be validated by marginal coverage alone.

2.5 Multivariate pre-rank calibration and regularization

Pre-rank functions reduce multivariate forecast–outcome pairs to scalar diagnostics whose PIT-like distribution can be checked for uniformity. Allen, Ziegel, and Ginsbourger develop this perspective for multivariate forecast assessment [18]. More directly, recent multi-output regression work develops differentiable pre-rank regularization for training multivariate predictive distributions [4, 19]. Those formulations include smoothed empirical PIT-CDF penalties; neither projection nor smoothing is claimed here as a new general multivariate calibration device. CA-RNN uses fixed linear projections as a narrower set of pre-ranks. Its adaptation concerns history-dependent recurrent forecasting on ordered data, the additional lagged PIT-moment

penalty, and evaluation under misspecification and distribution shift. Finite pooled restrictions still do not establish calibration conditional on every forecast history.

2.6 Probabilistic recurrent forecasting, misspecification, and shift

DeepAR demonstrates that autoregressive recurrent networks can estimate conditional predictive distributions by likelihood and share information across related series [20]. CA-RNN adopts the recurrent conditional-density perspective but uses a joint multivariate Gaussian head, then augments likelihood-based estimation with calibration restrictions. This separation is consequential under misspecification. White [21] shows that maximum likelihood estimation under a misspecified family generally converges to a pseudo-true distribution rather than the data-generating law. Calibration-aware regularization can move the fitted distribution away from that likelihood optimum, but such movement need not improve a strictly proper score. The exact regret arguments developed below formalize this tradeoff.

Distribution shift creates a second challenge. Ovadia et al. [22] show empirically that predictive uncertainty methods that appear adequate in distribution can deteriorate as dataset shift intensifies. Financial returns add pronounced conditional heteroskedasticity; the GARCH model of Bollerslev [23] provides the canonical likelihood-based mechanism for persistent time-varying variance. These strands motivate the manuscript’s shifted-regime experiments and its VAR-GARCH benchmark. A neural calibration method should not receive credit merely for learning volatility that a standard econometric model already represents directly.

2.7 Positioning of the present study

The literature already supplies differentiable calibration regularizers for continuous regression and multivariate pre-rank methods, as well as recurrent probabilistic prediction and PIT-based forecast diagnostics. The present study asks whether adapting a projected-PIT penalty to sequential recurrent forecasting, with an added temporal moment term, improves joint forecasts when the model is misspecified or the data regime changes. The primary contrast is the likelihood-trained RNN against CA-RNN; Bayesian variants support the methodological comparison. The confirmatory design tests whether gains in the optimized projected criterion transfer to proper scores, interval quality, tail-risk estimates, and an external asset panel. Its mixed findings delimit this adaptation rather than establishing a new general multivariate calibration method.

Table [A.1](#) makes the novelty boundary explicit. The distinction is not that recurrent networks alone make calibration sequential; it is that the loss and evaluation are applied to a filtration-indexed sequence of conditional forecasts and add finite lagged-PIT restrictions. These restrictions remain weaker than full conditional calibration.

Chapter 3

Forecasting Setup and CA–RNN Methodology

3.1 Forecasting setup and calibration

3.1.1 Sequential prediction problem

Let $(Y_t)_{t=1}^T$ be an adapted stochastic process with $Y_t \in \mathbb{R}^d$ and filtration $\mathcal{F}_t = \sigma(Y_1, \dots, Y_t)$. Given a history length H , define $X_t = (Y_{t-H+1}, \dots, Y_t)$. The task is to estimate a conditional predictive distribution

$$P_t(\cdot) = \mathbb{P}(Y_{t+1} \in \cdot \mid \mathcal{F}_t)$$

with a model $P_{\theta,t}$. Training, validation, and test blocks are ordered chronologically. Standardization parameters are estimated on the training block only.

3.1.2 Joint Gaussian predictive head

For the initial specification, the predictive law is

$$Y_{t+1} \mid \mathcal{F}_t, \theta \sim \mathcal{N}_d(\mu_{\theta,t}, \Sigma_{\theta,t}), \quad \Sigma_{\theta,t} = L_{\theta,t} L_{\theta,t}^\top. \quad (3.1)$$

The network outputs the d components of $\mu_{\theta,t}$ and the $d(d+1)/2$ lower-triangular entries of $L_{\theta,t}$. An unconstrained diagonal output $s_{t,j}$ is transformed as

$$L_{t,jj} = \text{softplus}(s_{t,j}) + \varepsilon, \quad \varepsilon > 0.$$

Proposition 3.1 (Validity of the predictive covariance). For every network output, $\Sigma_{\theta,t} = L_{\theta,t}L_{\theta,t}^\top$ is symmetric positive definite.

Proof. Symmetry follows from transposition. Since every diagonal entry of the triangular matrix $L_{\theta,t}$ is strictly positive, $L_{\theta,t}$ is nonsingular. Hence, for any nonzero $x \in \mathbb{R}^d$,

$$x^\top \Sigma_{\theta,t} x = x^\top L_{\theta,t} L_{\theta,t}^\top x = \|L_{\theta,t}^\top x\|_2^2 > 0.$$

Thus $\Sigma_{\theta,t}$ is positive definite. □

3.1.3 Negative log likelihood

For an index set I , the mean joint Gaussian negative log likelihood is

$$\mathcal{L}_{\text{NLL}}(\theta; I) = -\frac{1}{|I|} \sum_{t \in I} \log p_\theta(Y_{t+1} \mid X_t) \tag{3.2}$$

$$= \frac{1}{2|I|} \sum_{t \in I} \left[d \log(2\pi) + 2 \sum_{j=1}^d \log L_{t,jj} + e_t^\top \Sigma_{\theta,t}^{-1} e_t \right], \tag{3.3}$$

where $e_t = Y_{t+1} - \mu_{\theta,t}$. Computation uses triangular solves rather than an explicit matrix inverse.

Proposition 3.2 (Strict propriety of the population log score). Let $P_0(\cdot \mid \mathcal{F}_t)$ have

density p_0 and let p be any candidate conditional density with common support. Then

$$\mathbb{E}_{P_0}[-\log p(Y_{t+1} \mid \mathcal{F}_t)] - \mathbb{E}_{P_0}[-\log p_0(Y_{t+1} \mid \mathcal{F}_t)] = \mathbb{E}[\text{KL}(P_0(\cdot \mid \mathcal{F}_t) \| P(\cdot \mid \mathcal{F}_t))] \geq 0.$$

Equality holds only when the conditional densities agree almost surely.

Proof. Condition on $\mathcal{F}_t$, subtract the two expected log scores, and recognize the conditional KL divergence. Nonnegativity follows from Gibbs' inequality. Taking expectations over $\mathcal{F}_t$ proves the result. $\square$

This proposition explains why the calibration penalty must supplement rather than replace a proper score: calibration alone can admit uninformative forecasts.

3.1.4 Projected probability integral transforms

Fix unit vectors $a_1, \dots, a_R \in \mathbb{R}^d$. In the experiments they comprise the d coordinate directions, one equal-weight direction, and four random unit directions fixed by a documented seed. Define

$$U_{t,r} = F_{\theta,t,r}(a_r^\top Y_{t+1}), \tag{3.4}$$

where $F_{\theta,t,r}$ is the predictive CDF of $a_r^\top Y_{t+1}$.

Proposition 3.3 (Projected Gaussian law). Under (3.1),

$$a_r^\top Y_{t+1} \mid \mathcal{F}_t, \theta \sim \mathcal{N}(a_r^\top \mu_{\theta,t}, a_r^\top \Sigma_{\theta,t} a_r),$$

and therefore

$$U_{t,r} = \Phi\left(\frac{a_r^\top(Y_{t+1} - \mu_{\theta,t})}{\sqrt{a_r^\top \Sigma_{\theta,t} a_r}}\right).$$

Proof. A linear transformation of a multivariate Gaussian is Gaussian. Its mean and variance follow from linearity of expectation and the covariance transformation rule. Applying its CDF gives the displayed PIT. $\square$

Theorem 3.4 (Sequential PIT property). Suppose $F_{0,t,r}$ is the true, continuous conditional CDF of $a_r^\top Y_{t+1}$ given $\mathcal{F}_t$. Then $U_{t,r} = F_{0,t,r}(a_r^\top Y_{t+1})$ is uniformly distributed conditional on $\mathcal{F}_t$. For a fixed r , the sequence $(U_{t,r})_t$ is independent and $\text{Uniform}(0, 1)$.

Proof. For $u \in [0, 1]$, continuity and the conditional probability integral transform give

$$\mathbb{P}(U_{t,r} \leq u \mid \mathcal{F}_t) = u.$$

Thus $U_{t,r}$ is conditionally uniform and its conditional law does not depend on $\mathcal{F}_t$. For times $t_1 < \dots < t_m$, iterated conditioning yields

$$\begin{aligned} \mathbb{P}(U_{t_1,r} \leq u_1, \dots, U_{t_m,r} \leq u_m) &= \mathbb{E}\left[\mathbf{1}_{\{U_{t_1,r} \leq u_1, \dots, U_{t_{m-1},r} \leq u_{m-1}\}} \mathbb{P}(U_{t_m,r} \leq u_m \mid \mathcal{F}_{t_m})\right] \\ &= u_m \mathbb{P}(U_{t_1,r} \leq u_1, \dots, U_{t_{m-1},r} \leq u_{m-1}). \end{aligned}$$

Induction gives $\prod_{j=1}^m u_j$, proving independence and uniformity. $\square$

Remark 3.1. Uniformity for finitely many a_r is necessary but not sufficient for equality of multivariate distributions. If all one-dimensional projections agree, the Cramér–Wold device identifies the joint distribution; a finite projection set only provides a tractable approximation and diagnostic.

Training uses the analytic Gaussian PIT above. The frozen external-panel out-of-sample evaluation uses the common sample-rank approximation

$$\hat{U}_{t,r} = \frac{\sum_{s=1}^S \mathbf{1}\{a_r^\top Y_{t+1}^{(s)} < a_r^\top Y_{t+1}\} + 1/2}{S + 1},$$

with $S = 500$ predictive draws. This supports the same diagnostic for deterministic Gaussian forecasts and Bayesian posterior-predictive mixtures, at the cost of finite-simulation error. The post-protocol evaluation audit additionally computes exact projected PITs for deterministic Gaussian neural forecasts; these are reported separately rather than substituted for the frozen sample-rank outcome.

3.2 Calibration-aware RNN methodology

3.2.1 Recurrent architecture

For a single-layer GRU, with elementwise product $\odot$,

$$\begin{aligned} r_t &= \sigma(W_{ir}Y_t + b_{ir} + W_{hr}h_{t-1} + b_{hr}), \\ z_t &= \sigma(W_{iz}Y_t + b_{iz} + W_{hz}h_{t-1} + b_{hz}), \\ n_t &= \tanh(W_{in}Y_t + b_{in} + r_t \odot (W_{hn}h_{t-1} + b_{hn})), \\ h_t &= n_t + z_t \odot (h_{t-1} - n_t). \end{aligned}$$

Linear heads map the final hidden state to $\mu_{\theta,t}$ and the unconstrained entries of $L_{\theta,t}$.

3.2.2 Differentiable projected calibration loss

Let $q_g = g/(G + 1)$ for $g = 1, \dots, G$. The empirical CDF of PIT values is nondifferentiable in model parameters because it uses indicators. Replace the indicator $\mathbf{1}\{U_{t,r} \leq q_g\}$ by

$$s_\tau(q_g - U_{t,r}) = \sigma\left(\frac{q_g - U_{t,r}}{\tau}\right), \quad \tau > 0.$$

For a consecutive training batch B , define

$$\mathcal{L}_{\text{cal}}(\theta; B) = \frac{1}{GR} \sum_{g=1}^G \sum_{r=1}^R \left[\frac{1}{|B|} \sum_{t \in B} s_\tau(q_g - U_{t,r}) - q_g \right]^2. \quad (3.5)$$

This is a grid approximation to an integrated squared CDF discrepancy. Its smoothed empirical PIT-CDF form belongs to the differentiable regression and multivariate pre-rank regularization family [3, 4, 19]; it is analogous in purpose, but not in definition, to the classification AECE in [1]. The present instantiation uses fixed linear projections and analytic Gaussian CDFs within a recurrent conditional forecast, rather than proposing the discrepancy as a new general-purpose penalty.

Proposition 3.5 (Differentiability). If the network outputs are differentiable in θ and $a_r^\top \Sigma_{\theta,t} a_r > 0$, then $\mathcal{L}_{\text{cal}}(\theta; B)$ is differentiable in θ .

Proof. The Cholesky parameterization, affine projection, positive square root, Gaussian CDF, logistic smoother, finite sums, and squaring operations are differentiable on the stated domain. Their composition is therefore differentiable. $\square$

Proposition 3.6 (Zero-temperature limit). Fix PIT values such that $U_{t,r} \neq q_g$ for every t, r, g . As $\tau \downarrow 0$, (3.5) converges to the squared grid discrepancy formed with the empirical PIT CDF.

Proof. For nonzero x , $\sigma(x/\tau) \rightarrow \mathbf{1}\{x > 0\}$ as $\tau \downarrow 0$. Every sum is finite, so the limit passes through the averages and squares. The exclusion of ties removes the ambiguous boundary value. $\square$

Proposition 3.7 (Finite-grid interpretation). For fixed B , projections, grid, and $\tau > 0$, $\mathcal{L}_{\text{cal}}(\theta; B) = 0$ if and only if

$$\frac{1}{|B|} \sum_{t \in B} s_\tau(q_g - U_{t,r}) = q_g \quad \text{for every } (g, r).$$

Consequently, zero training loss establishes agreement only at the selected smoothed CDF constraints; it does not establish uniformity between grid points, for untested projections, conditional on the history, or out of sample.

Proof. Equation (3.5) is an average of nonnegative squared discrepancies. It is zero exactly when each squared discrepancy is zero. The stated limitations follow because the objective contains no other grid values, projections, conditioning events, or observations outside B . $\square$

3.2.3 Temporal PIT-moment loss

PIT uniformity does not rule out forecastable temporal structure. With $\tilde{U}_{t,r} = U_{t,r} - 1/2$ and maximum lag K , define

$$\mathcal{L}_{\text{seq}}(\theta; B) = \frac{1}{RK} \sum_{r=1}^R \sum_{k=1}^K \left[\frac{1}{|B| - k} \sum_{t \in B: k \text{ available}} \tilde{U}_{t,r} \tilde{U}_{t-k,r} \right]^2. \quad (3.6)$$

Under the ideal conditional forecast, the population quantities targeted by (3.6) are zero by the sequential PIT theorem. The converse is not true: finitely many zero autocovariances do not establish independence.

Proposition 3.8 (Exact scope of the temporal PIT-moment loss). Suppose $|B| > K$. Then $\mathcal{L}_{\text{seq}}(\theta; B) = 0$ if and only if

$$\frac{1}{|B| - k} \sum_{t \in B: k \text{ available}} (U_{t,r} - 1/2)(U_{t-k,r} - 1/2) = 0 \quad \text{for all } r = 1, \dots, R, \ k = 1, \dots, K.$$

If the forecast is correctly specified so that each projected sequential PIT is i.i.d. $\text{Unif}(0, 1)$, every corresponding population moment is zero. Conversely, satisfaction of these finitely many moments does not imply temporal independence, conditional uniformity, or correct joint dynamics.

Proof. The first statement again follows because (3.6) is a sum of squares. Under i.i.d. uniform PITs, $\mathbb{E}(U_{t,r} - 1/2) = 0$ and independence gives $\mathbb{E}[(U_{t,r} - 1/2)(U_{t-k,r} - 1/2)] = 0$ for $k \geq 1$. The converse fails because a finite set of second-order moment restrictions cannot characterize an entire serial joint law; dependent processes can be uncorrelated at the tested lags. $\square$

Remark 3.2. The centered product in (3.6) uses the uniform-law mean $1/2$, not the batch PIT mean. It is therefore a calibration moment rather than the usual sample autocovariance unless the batch mean is exactly $1/2$.

3.2.4 Frequentist CA-RNN objective

The proposed frequentist objective is

$$\hat{\theta}_{\text{CA}} = \arg \min_{\theta} \{ \mathcal{L}_{\text{NLL}}(\theta; \mathcal{D}) + \lambda_{\text{cal}} \mathcal{L}_{\text{cal}}(\theta; \mathcal{D}) + \lambda_{\text{seq}} \mathcal{L}_{\text{seq}}(\theta; \mathcal{D}) \}. \quad (3.7)$$

Both calibration terms in (3.7) define CA-RNN. The restriction $\lambda_{\text{seq}} = 0$ is used only as a loss-component ablation that removes the temporal PIT-moment penalty.

3.2.5 Supporting Bayesian extension

The Bayesian construction is retained to mirror the CA-BNN comparison and test whether parameter averaging changes the calibration result. It is not the primary methodological or empirical claim of this paper; the main estimand is the frequentist RNN-versus-CA-RNN contrast.

Let every scalar neural parameter have variational posterior and prior

$$q_{\phi}(\theta_j) = \mathcal{N}(\mu_j, \sigma_j^2), \quad p(\theta_j) = \mathcal{N}(0, \sigma_0^2), \quad \sigma_j = \text{softplus}(\rho_j) + \epsilon.$$

Sampling uses the reparameterization

$$\theta_j = \mu_j + \sigma_j \varepsilon_j, \quad \varepsilon_j \sim \mathcal{N}(0, 1).$$

The independent-Gaussian KL term has closed form

$$\text{KL}(q_{\phi} \| p) = \sum_j \left[\log \frac{\sigma_0}{\sigma_j} + \frac{\sigma_j^2 + \mu_j^2}{2\sigma_0^2} - \frac{1}{2} \right]. \quad (3.8)$$

The calibration-aware variational objective is

$$\mathcal{F}^{\text{CA}}(\phi) = \mathbb{E}_{\theta \sim q_\phi} [\mathcal{L}_{\text{NLL}}(\theta; \mathcal{D}) + \lambda_{\text{cal}} \mathcal{L}_{\text{cal}}(\theta; \mathcal{D}) + \lambda_{\text{seq}} \mathcal{L}_{\text{seq}}(\theta; \mathcal{D})] \quad (3.9)$$

$$+ \frac{\beta}{N} \text{KL}(q_\phi(\theta) \| p(\theta)). \quad (3.10)$$

This is the continuous-outcome recurrent analogue of the CA-BNN free energy in [1].

The factor $1/N$ matches the mean per-observation likelihood convention. The weight β remains a validation-selected generalized-Bayes hyperparameter.

Proposition 3.9 (Unbiased reparameterized gradient for the expectation). Assume differentiation and expectation may be interchanged. If $\theta^{(m)} = \mu + \sigma(\rho) \odot \varepsilon^{(m)}$ with independent $\varepsilon^{(m)} \sim \mathcal{N}(0, I)$, then the Monte Carlo gradient of the expected data-dependent part of (3.10) is unbiased.

Proof. Write the expectation with respect to the parameter-free distribution of ε . Under the interchange assumption,

$$\nabla_\phi \mathbb{E}_\varepsilon [f(\mu + \sigma(\rho) \odot \varepsilon)] = \mathbb{E}_\varepsilon \nabla_\phi f(\mu + \sigma(\rho) \odot \varepsilon).$$

The sample average of independent terms has this expectation. The KL gradient is computed analytically from (3.8). $\square$

Remark 3.3 (Dependent mini-batches). Unlike independent classification observations, overlapping time-series windows are dependent. Consecutive batches are required to evaluate $\mathcal{L}_{\text{seq}}$, but their nonlinear calibration loss is not asserted to be an unbiased estimate of the full-sample calibration loss. Batch length and history overlap are therefore methodological hyperparameters and sensitivity targets.

In the implementation, both calibration penalties are batch-local: lag pairs crossing mini-batch boundaries are omitted. Epoch summaries average batch means, so the final shorter batch receives the same summary weight as a full batch. The displayed dataset-level objectives are the target criteria; stochastic optimization uses this documented batch-local surrogate.

3.2.6 Why targeted calibration need not improve proper scores

This section gives a theoretical interpretation of the negative confirmatory result. The central distinction is between satisfying a finite collection of calibration moments and estimating the complete conditional distribution. The former is the explicit target of CA-RNN; the latter is assessed by strictly proper scores.

3.2.6.1 Population target and finite-projection non-identification

For a fixed forecast law Q and projection a_r , let $H_{Q,r}(u) = \mathbb{P}_Q(U_{t,r} \leq u)$ denote the CDF of the resulting PIT under the true data law. The unsmoothed population counterpart of (3.5) on the implemented grid is

$$\mathcal{C}_{R,G}(Q) = \frac{1}{RG} \sum_{r=1}^R \sum_{g=1}^G \{H_{Q,r}(q_g) - q_g\}^2. \quad (3.11)$$

It measures unconditional agreement at RG points. It is not a metric on the space of joint conditional distributions.

Proposition 3.10 (Finite-grid calibration is non-identifying). Even in one dimension with a continuous strictly increasing true CDF F , there are infinitely many forecast laws $Q \neq P$ for which $\mathcal{C}_{1,G}(Q) = 0$ on any fixed finite grid.

Proof. Choose any strictly increasing bijection $h : [0, 1] \rightarrow [0, 1]$ that fixes every grid point, $h(q_g) = q_g$, but is not the identity between at least two adjacent grid points. Define forecast law Q by its CDF $G(y) = h(F(y))$. If $Y \sim F$, then $G(Y) = h(F(Y))$ and, because $F(Y)$ is uniform,

$$\mathbb{P}\{G(Y) \leq q_g\} = h^{-1}(q_g) = q_g.$$

Thus every grid discrepancy is zero, while $G \neq F$. Infinitely many such monotone deformations h can be constructed between grid points. $\square$

The multivariate case has two additional sources of non-identification: only a finite set of directions is inspected, and the moments are averaged across those directions. The Cramér–Wold device would identify a joint law if all one-dimensional projected distributions agreed, but it does not turn a finite projection set into an identifying criterion.

Proposition 3.11 (Consistency of fixed calibration moments). Fix Q , R , G , and $\tau > 0$. If the PIT process is stationary and ergodic, then the empirical smoothed grid moments converge almost surely to their population counterparts. The empirical lagged moments in (3.6) likewise converge almost surely for every fixed tested lag, provided their products are integrable.

Proof. The smoothed indicators are bounded measurable functions of the PIT process, so Birkhoff’s ergodic theorem applies to each sample average. The lagged products are measurable functions of the finite-dimensional shifted process and are integrable by assumption. There are finitely many projections, grid points, and lags, so convergence

holds jointly for all implemented moments. $\square$

This result concerns a fixed forecast rule under one stationary law. It does not provide uniform convergence over neural parameters, correct batch-local optimization bias, or transfer the moments to a shifted test law. Those missing implications are precisely relevant to the observed training-to-test gap.

3.2.6.2 Zero PIT autocovariance does not imply independence

The limitation of the sequential penalty is constructive, not merely technical.

Proposition 3.12 (Dependent uniform PITs with zero covariance). There exist dependent $U, V \sim \text{Unif}(0, 1)$ satisfying $\mathbb{E}[(U - 1/2)(V - 1/2)] = 0$.

Proof. Let $f(u) = 6u^2 - 6u + 1$ and define a bivariate density on $[0, 1]^2$ by

$$c_\varepsilon(u, v) = 1 + \varepsilon f(u)f(v), \quad 0 < |\varepsilon| < 1.$$

Because $\int_0^1 f(u) du = 0$, both marginals are uniform. The chosen range of ε keeps the density nonnegative. Since c_ε is not constant, U and V are dependent. However, symmetry of f about $1/2$ gives

$$\int_0^1 (u - 1/2)f(u) du = 0,$$

and hence

$$\mathbb{E}[(U - 1/2)(V - 1/2)] = \varepsilon \left\{ \int_0^1 (u - 1/2)f(u) du \right\}^2 = 0.$$

□

Consequently, even population minimization of every moment represented by a finite-lag version of (3.6) cannot establish PIT independence. The confirmatory increase in PIT autocorrelation is stronger still: it shows that the estimated training moments did not generalize as lower corresponding test moments.

3.2.6.3 Energy Score regret is not controlled by calibration loss

For distributions with finite first moments, define

$$\text{ES}(Q, y) = \mathbb{E}\|X - y\|_2 - \frac{1}{2}\mathbb{E}\|X - X'\|_2, \quad X, X' \stackrel{\text{iid}}{\sim} Q.$$

Proposition 3.13 (Energy Score regret). If $Y \sim P$, then

$$\mathbb{E}_P \text{ES}(Q, Y) - \mathbb{E}_P \text{ES}(P, Y) = \frac{1}{2}\mathcal{E}(P, Q) \geq 0, \quad (3.12)$$

where

$$\mathcal{E}(P, Q) = 2\mathbb{E}\|X - Y\|_2 - \mathbb{E}\|X - X'\|_2 - \mathbb{E}\|Y - Y'\|_2$$

is the energy distance. Equality holds only when $P = Q$.

Proof. Expand the two expected Energy Scores using independent $X, X' \sim Q$ and $Y, Y' \sim P$, then subtract. The remaining expression is one half of the energy distance. Its nonnegativity and identity of indiscernibles are the defining properties of energy distance on distributions with finite first moments. □

Equation (3.12) shows that Energy Score evaluates the complete joint distribution, whereas (3.11) evaluates only finitely many scalar constraints. No inequality in the CA-RNN objective bounds $\mathcal{E}(P, Q)$ by $\mathcal{C}_{R,G}(Q)$.

Proposition 3.14 (Calibration improvement need not improve a strictly proper score). For any fixed finite PIT grid, there exist a true continuous distribution P and forecasts Q_0, Q_1 such that

$$\mathcal{C}_{1,G}(Q_1) < \mathcal{C}_{1,G}(Q_0), \quad \mathbb{E}_P S(Q_1, Y) > \mathbb{E}_P S(Q_0, Y)$$

for any strictly proper score S that is finite and continuous along sufficiently small CDF perturbations of P .

Proof. Use the preceding monotone construction to choose $Q_1 \neq P$ with zero finite-grid calibration loss. Strict propriety gives a positive score regret for Q_1 . Now choose $Q_0 \neq P$ as an arbitrarily small monotone perturbation of P that does not fix at least one grid point. Then $\mathcal{C}_{1,G}(Q_0) > 0$, while continuity makes its proper-score regret arbitrarily close to zero. Choose the perturbation small enough that its regret is below that of Q_1 . The two displayed inequalities follow. $\square$

This proposition formalizes the main empirical result: a lower projected-PIT error is not an ordering of forecast distributions and cannot, by itself, imply a lower Energy Score.

3.2.6.4 *Why interval score can also worsen*

Let F be the true marginal CDF, α the nominal miscoverage probability, $\tau = \alpha/2$, and $(\ell^*, u^*) = (F^{-1}(\tau), F^{-1}(1 - \tau))$ the optimal central interval. For a candidate interval (ℓ, u) , the expected interval-score regret has an exact quantile representation.

Proposition 3.15 (Interval-score regret). For continuous F with finite first moment,

$$\mathbb{E}_F \text{IS}_\alpha(\ell, u; Y) - \mathbb{E}_F \text{IS}_\alpha(\ell^*, u^*; Y) \quad (3.13)$$

$$= \frac{1}{\tau} \left\{ \int_{\ell^*}^{\ell} [F(x) - \tau] dx + \int_{u^*}^u [F(x) - (1 - \tau)] dx \right\} \geq 0. \quad (3.14)$$

Proof. The interval score equals τ^{-1} times the sum of pinball losses at levels τ and $1 - \tau$. The derivative of expected pinball loss with respect to a candidate quantile q is $F(q) - p$ at level p . Integrating this derivative from each true quantile to its candidate value yields (3.14). Monotonicity of F makes both oriented integrals nonnegative. $\square$

The CA-RNN loss averages 19 grid levels and multiple directions and is optimized with smoothing on dependent training batches. A reduction in that average need not reduce coordinate-specific errors at the 0.05 and 0.95 quantiles on a shifted test distribution. Equation (3.14) shows that any such tail-quantile displacement creates positive expected interval-score regret. This explains how projected calibration can improve while the interval score worsens: the scalar calibration average falls, but errors at the particular marginal tail functionals priced by the interval score increase or the added width is not offset by sufficiently smaller miss penalties.

Algorithm 1 Frequentist CA-RNN training

Require: Ordered observations $\mathcal{D}$; history H ; projections A ; weights $\lambda_{\text{cal}}, \lambda_{\text{seq}}$; epoch limit E ; gradient threshold c

Ensure: Fitted parameters $\hat{\theta}$ selected without test information

- 1: Split $\mathcal{D}$ chronologically into $\mathcal{D}_{\text{tr}}, \mathcal{D}_{\text{val}}, \mathcal{D}_{\text{te}}$
- 2: Fit standardization on $\mathcal{D}_{\text{tr}}$ and construct length- H windows
- 3: Initialize θ and fix $A = (a_1, \dots, a_R)^\top$
- 4: **for** $e = 1, \dots, E$ **do**
- 5: **for** consecutive mini-batch $B \subset \mathcal{D}_{\text{tr}}$ **do**
- 6: $(\mu_{\theta,t}, L_{\theta,t})_{t \in B} \leftarrow \text{CARNNFORWARD}(B, \theta)$
- 7: Compute $\mathcal{L}_{\text{NLL}}$ by (3.3) and projected PITs by (3.4)
- 8: Compute $\mathcal{L}_{\text{cal}}$ by (3.5) and $\mathcal{L}_{\text{seq}}$ by (3.6)
- 9: $\mathcal{J} \leftarrow \mathcal{L}_{\text{NLL}} + \lambda_{\text{cal}}\mathcal{L}_{\text{cal}} + \lambda_{\text{seq}}\mathcal{L}_{\text{seq}}$
- 10: $g \leftarrow \text{CLIPGRADNORM}(\nabla_{\theta}\mathcal{J}, c)$
- 11: Update θ with AdamW using g
- 12: **end for**
- 13: Evaluate the prespecified selection score on $\mathcal{D}_{\text{val}}$
- 14: Save θ if the validation score improves
- 15: **if** the patience criterion is met **then**
- 16: **break**
- 17: **end if**
- 18: **end for**
- 19: Restore the selected checkpoint $\hat{\theta}$
- 20: Evaluate $\hat{\theta}$ once on $\mathcal{D}_{\text{te}}$

3.2.7 Training algorithms

For fixed-budget causal loss ablations, checkpoint selection and the patience test in Algorithm 1 are disabled, and all specifications are compared after the same number of epochs.

Algorithm 2 is the recurrent continuous-outcome adaptation of the CA-BNN procedure in [1].

Algorithm 2 Variational calibration-aware Bayesian RNN training

Require: Ordered observations $\mathcal{D}$; projections A ; prior scale σ_0 ; Monte Carlo count M ; loss weights $\lambda_{\text{cal}}, \lambda_{\text{seq}}, \beta$; epoch limit E

Ensure: Variational parameters $\hat{\phi} = (\hat{\mu}, \hat{\rho})$

- 1: Construct chronological splits and training-only standardization as in Algorithm 1
- 2: Initialize $\phi = (\mu, \rho)$ and set $\sigma(\rho) = \text{softplus}(\rho) + \epsilon$
- 3: **for** $e = 1, \dots, E$ **do**
- 4: **for** consecutive mini-batch $B \subset \mathcal{D}_{\text{tr}}$ **do**
- 5: **for** $m = 1, \dots, M$ **do**
- 6: Draw $\varepsilon^{(m)} \sim \mathcal{N}(0, I)$ and set $\theta^{(m)} \leftarrow \mu + \sigma(\rho) \odot \varepsilon^{(m)}$
- 7: Use $\theta^{(m)}$ consistently at every recurrent time step
- 8: Compute $\mathcal{J}_{\text{data}}^{(m)} \leftarrow \mathcal{L}_{\text{NLL}}^{(m)} + \lambda_{\text{cal}} \mathcal{L}_{\text{cal}}^{(m)} + \lambda_{\text{seq}} \mathcal{L}_{\text{seq}}^{(m)}$
- 9: **end for**
- 10: $\hat{\mathcal{F}}^{\text{CA}} \leftarrow M^{-1} \sum_{m=1}^M \mathcal{J}_{\text{data}}^{(m)} + (\beta/N) \text{KL}(q_{\phi} \| p)$
- 11: Differentiate $\hat{\mathcal{F}}^{\text{CA}}$, clip its gradient, and update ϕ with AdamW
- 12: **end for**
- 13: Evaluate validation NLL at posterior-mean weights and update the selected checkpoint
- 14: **if** the patience criterion is met **then**
- 15: **break**
- 16: **end if**
- 17: **end for**
- 18: Restore $\hat{\phi}$
- 19: **for** each test origin t and predictive draw s **do**
- 20: Draw $\theta^{(s)} \sim q_{\hat{\phi}}$ and then $Y_{t+1}^{(s)} \sim p_{\theta^{(s)}}(\cdot | \mathcal{F}_t)$
- 21: **end for**
- 22: Evaluate the posterior predictive mixture on $\mathcal{D}_{\text{te}}$

Chapter 4

Distribution Shift and Experimental Design

4.1 Model misspecification and distribution shift

4.1.1 Distinguishing misspecification from shift

Model misspecification occurs when the true conditional law $P_0(Y_{t+1} \mid \mathcal{F}_t)$ is outside the fitted family, even if the data law is stationary. Distribution shift occurs when the evaluation law differs from the development law. The two can coexist: a Gaussian predictive head is misspecified under heavy-tailed innovations, while a transition from Gaussian training innovations to heavy-tailed test innovations is also a distribution shift.

Let P_0 denote the reference law and let $P_{\delta,s}$ denote stress family s at severity $\delta \geq 0$, with $P_{0,s} = P_0$. For metric M , define degradation

$$\Delta_M^{(m,s)}(\delta) = M(P_{\delta,s}, \hat{P}_m) - M(P_{0,s}, \hat{P}_m), \quad (4.1)$$

where m indexes the forecasting method. For metrics in which smaller is better, positive Δ_M denotes deterioration. Coverage is reported as absolute error from the nominal level before differencing.

4.1.2 Stress families

The following stress families distinguish mechanisms; only the first has been executed as a replicated isolated-factor experiment, and none has been executed as a full factorial study:

1. **Scale shift:** innovation covariance is multiplied by a severity function while dependence is held fixed.
2. **Dependence shift:** correlations or a dynamic correlation state change while marginal variances are controlled.
3. **Tail shift:** Gaussian innovations transition to Student- t , mixtures, or asymmetric downside contamination.
4. **Volatility misspecification:** ARCH/GARCH or stochastic-volatility dynamics violate the Gaussian head's conditionally homoskedastic approximation.
5. **Mean-dynamics shift:** nonlinear interactions or an abrupt change in recurrent coefficients alter the conditional location.

The initial developmental curve uses common random numbers and jointly changes scale and common correlation. Models are selected without examining its stressed test metrics and are frozen after fitting at $\delta = 0$. The completed post-protocol experiment isolates scale: 50 independent reference datasets are generated, paired RNN and CA-RNN models are fitted on each unshifted training history, and the same fitted models are evaluated at scale severities $\delta \in \{0, 0.5, 1\}$. The innovation

standard deviation becomes $1 + \delta$ times its reference value only in the test partition; correlation, transition coefficients, and innovation family remain fixed. Tail, volatility, and mean-dynamics regimes appear in separate developmental simulations, not in this isolated-factor experiment.

4.1.3 Robustness estimands

For paired realization b , method m , comparator c , and metric M , define

$$D_b^{m,c}(\delta) = M_b^m(\delta) - M_b^c(\delta).$$

For the completed isolated-scale experiment with $B = 50$ independent DGP realizations, the summary is

$$\overline{D}^{m,c}(\delta) = \frac{1}{B} \sum_{b=1}^B D_b^{m,c}(\delta), \quad \text{MCSE}(\overline{D}) = \frac{s_D(\delta)}{\sqrt{B}}.$$

One distinct neural seed is assigned to each DGP realization; thus the reported MCSE includes both DGP and optimization variation but does not identify their separate contributions. Failure frequencies and runtimes are reported rather than discarded implementation details. The comparison is developmental because the model weights and shift family were selected after earlier exploratory analyses; it is not an additional frozen confirmatory replication.

4.2 Experimental design

4.2.1 Primary comparison and supporting Bayesian replication

The design follows the comparison logic of [1], but the first and third rows define the primary frequentist comparison. The Bayesian rows are supporting replications:

Architectures, histories, splits, optimizers, stopping rules, and predictive heads are matched wherever the Bayesian formulation permits. The same projection matrix is used for all models in a comparison.

4.2.2 Conventional benchmark ladder

The neural comparisons are supplemented by three increasingly dynamic probabilistic benchmarks. A centered historical bootstrap samples joint training returns and contains neither a dynamic conditional mean nor time-varying scale. A Gaussian VAR estimates a linear conditional mean and a fixed innovation covariance. The strongest conventional benchmark combines that VAR mean with separate marginal GARCH(1, 1) filters,

$$h_{j,t} = \omega_j + \alpha_j e_{j,t-1}^2 + \beta_j h_{j,t-1}, \quad \omega_j > 0, \quad \alpha_j, \beta_j \geq 0, \quad \alpha_j + \beta_j < 1,$$

and samples entire standardized-residual vectors $z_t = (e_{1,t}/\sqrt{h_{1,t}}, \dots, e_{d,t}/\sqrt{h_{d,t}})$ from the training period. Its one-step samples are

$$Y_t^{(s)} = \hat{\mu}_t + \text{diag}(\sqrt{\hat{h}_t}) z_{I_s},$$

where I_s is a common row index across margins. This preserves empirical contemporaneous dependence and tail shape while allowing marginal conditional volatility to evolve. The VAR lag is selected from $\{1, 5, 10, 20\}$ by validation one-step MSE; parameters and the residual resampling distribution use training data only.

Model checking for the selected benchmark is also training-only. VAR dynamic stability is assessed by the spectral radius of its companion matrix. GARCH admissibility is checked through optimizer convergence, nonnegative parameters, $\alpha_j + \beta_j < 1$, and the frequency with which filtered variances reach the numerical floor. Ljung–Box statistics at lags 5, 10, and 20 are computed for standardized residuals and their squares to diagnose remaining conditional-mean and volatility dependence. Standardized-residual skewness, kurtosis, Jarque–Bera tests, and the residual correlation matrix assess the rationale for resampling empirical residual vectors jointly. The Ljung–Box and Jarque–Bera p -values are descriptive because they do not account for VAR–GARCH parameter estimation or multiplicity.

4.2.3 Frozen external-panel replication

The confirmatory protocol is stored separately in `docs/confirmatory_protocol.md` before outcome inspection. It fixes a disjoint MSFT–BAC–CVX–COST panel, the date boundary, transformations, split, architecture, loss weights, training and forecast seeds, primary outcomes, baseline definitions, bootstrap procedure, and failure reporting. The sole primary contrast is CA-RNN minus RNN across ten matched optimization seeds; conventional-baseline comparisons are secondary. Any post-outcome change creates an exploratory analysis rather than replacing the frozen result. The archived

protocol retains the model label used when it was frozen; that label refers to the identical projected-plus-temporal objective now called CA-RNN.

4.2.4 Calibration-weight experiment

Analogously to the source paper’s plot of ECE and accuracy against λ , this study plots projected-PIT calibration error and predictive performance against λ_{cal} . The paired panels will contain:

1. projected calibration discrepancy and reliability curves;
2. Energy Score, NLL where comparable, RMSE, interval score, and width.

The result of interest is a Pareto relationship, not calibration alone. Separate curves are required for CA-RNN and CA-BRNN.

4.2.5 Temporal-penalty ablation

The minimum ablation is:

The Bayesian versions repeat this comparison with the KL term always present. Additional ablations vary coordinate-only versus augmented projections, smoothing temperature, batch length, lag order, prior scale, and β .

To prevent checkpoint selection from masquerading as a calibration effect, the primary loss-component ablation uses identical initial parameters, chronological mini-batches, optimization budgets, and predictive random numbers, with early stopping disabled. A secondary operational comparison may use early stopping, but all methods must

then be selected by the same validation NLL rule.

4.2.6 Controlled simulations

Development simulations include Gaussian, heavy-tailed, heteroskedastic, and nonlinear regimes. A future confirmatory experiment should use independently generated data for every regime, sample size, shift severity, and DGP seed, with the following fixed in advance:

- sample sizes and number of independent realizations;
- neural-seed subset and number of optimization repetitions;
- calibration and KL grids;
- projection directions and smoothing grid;
- primary metrics and failure definitions; and
- rules for early stopping and invalid numerical runs.

4.2.7 Financial application

The application uses adjusted daily prices for AAPL, JPM, XOM, and WMT beginning in 2007 and converts them to aligned log returns. The four assets provide a small, computationally manageable multivariate system with heterogeneous sector exposure. All splits are chronological. Development selected $\lambda_{\text{cal}} = 100$ and $\lambda_{\text{seq}} = 2000$ before the multi-seed and rolling-regime analyses. The focused frequentist comparison uses

a 32-unit GRU, a 20-day history, a joint Gaussian head, 500 forecast draws, and validation-NLL early stopping.

Standardization is fitted on each training partition. Validation contexts carry the final 20 training observations, and test contexts carry the final 20 validation observations; every evaluation target therefore retains an available historical context without using future information. RMSE, Energy Score, interval width, and interval score are reported in training-standardized units. Coverage and PIT diagnostics are invariant to the componentwise affine transformation. The two penalty weights were chosen after inspecting developmental simulation and ablation outcomes, rather than by a pristine validation-only or preregistered selection procedure.

The rolling analysis uses expanding estimation samples and four non-overlapping test periods: 2017–2019 (pre-pandemic), 2020–2021 (pandemic), 2022–2023 (inflation and monetary tightening), and 2024–2025 (recent). Each window has a strictly subsequent validation period and five matched optimization seeds. These labels describe calendar periods rather than uniquely identified causal regimes. Because the full panel was observed during development and the final two windows overlap the initial test period, this is temporal robustness evidence rather than a pristine preregistered confirmatory experiment.

4.2.8 Status of the evidence

The lambda sweep, ablation, original-panel ten-seed study, rolling windows, single-realization mixed-shift curve, and replicated isolated-scale experiment are developmental or internal-validation evidence. They condition on a specification informed

by earlier results. The external-panel replication is confirmatory for its single primary contrast: CA-RNN minus RNN on MSFT, BAC, CVX, and COST under the protocol frozen in `docs/confirmatory_protocol.md`. Its traditional-baseline comparisons were prespecified but secondary. Inferential intervals quantify test-origin and optimization-seed uncertainty under the implemented design; they do not include uncertainty over asset populations, architectures, or alternative research decisions.

4.2.9 Paired uncertainty design

For model m , comparator c , seed s , and forecast origin t , let $d_{s,t}^{m,c}$ be a paired additive-score difference. The estimand is

$$\hat{\Delta}^{m,c} = \frac{1}{S} \sum_{s=1}^S \frac{1}{T} \sum_{t=1}^T d_{s,t}^{m,c}.$$

The hierarchical bootstrap resamples optimization seeds and, within each selected seed, matched forecast origins in moving blocks of 20 trading days. Projected-PIT calibration, coverage error, and PIT dependence are recomputed in every bootstrap draw. The initial financial analysis uses 2,000 draws and each rolling regime uses 1,000. Raw coverage and width remain descriptive; the inferential coverage loss is

$$E_{\text{cov}} = |\hat{C}_{0.90} - 0.90|.$$

Negative candidate-minus-comparator differences favor the candidate for loss metrics. An interval containing zero is called unresolved, not evidence of equality. Inference conditions on the architecture, tuning path, projection set, assets, and observed calendar windows.

4.2.10 Metrics and figures

The tables report mean performance, paired differences, and uncertainty intervals.

The diagnostic figures include:

1. calibration and predictive metrics versus λ_{cal} ;
2. nominal-versus-empirical coverage curves;
3. projected PIT histograms or CDF reliability curves;
4. PIT autocorrelation diagnostics;
5. degradation curves versus shift severity; and
6. a calibration-sharpness or calibration-Energy Score frontier.

Chapter 5

Results and Conclusions

5.1 Results

5.1.1 Implementation validation

The implementation tests cover tensor dimensions, covariance positive-definiteness, loss evaluation, forecasting, chronological data handling, traditional baselines, and paired-bootstrap identities. Thirty-one automated tests pass. The notebooks also completed the simulation, shift, financial, optimization-seed, and rolling-regime workflows. Successful execution establishes implementation consistency but is not itself statistical validation.

5.1.2 Primary comparison and supporting Bayesian results

The refreshed matched nonlinear simulation showed that calibration-aware objectives alter the forecast distribution but do not yield uniform dominance. Relative to RNN, CA-RNN slightly reduced RMSE (0.9368 to 0.9334), Energy Score (1.2354 to 1.2348), and mean absolute projected-PIT autocorrelation (0.0733 to 0.0674). It increased interval score (3.8771 to 3.9297) and projected-PIT calibration error (0.001224 to 0.001455), despite improving nominal coverage from 0.9167 to 0.9375 only in the sense of producing wider intervals; both coverage rates exceed the 0.90 target. Bayesian averaging did not remove this calibration–score tradeoff: CA-BRNN

had worse RMSE, Energy Score, interval score, and projected calibration than BRNN, although its projected-PIT temporal dependence was slightly lower. RNN and CA-RNN therefore form the focused empirical comparison; BRNN and CA-BRNN remain supporting replications of the source paper’s model matrix. These are single-realization developmental results and are not treated as inferential evidence.

5.1.3 Calibration-weight sensitivity

The exploratory sweep over $\lambda_{\text{cal}} \in \{0, 1, 10, 50, 100, 250, 500\}$ displayed a Pareto tradeoff: projected calibration improved through the intermediate penalty range, whereas likelihood and proper-score performance deteriorated. The value 100 was retained as a developmental compromise rather than selected to maximize a final test result. Calibration improvement alone is therefore not interpreted as a better probabilistic forecast.

5.1.4 Temporal-penalty ablation

The temporal term has regime- and design-dependent effects. In the refreshed single-realization heteroskedastic simulation, removing the temporal penalty gave Energy Score 1.2101, interval score 3.9236, RMSE 0.9222, projected-PIT calibration error 0.000770, and mean absolute projected-PIT autocorrelation 0.05182. The default CA-RNN temporal weight of 2,000 instead gave 1.2153, 3.9409, 0.9258, 0.001046, and 0.05132, respectively. Thus the temporal term slightly reduced the diagnostic it directly targets, but worsened the proper scores, point error, projected uniformity, and absolute nominal-coverage error in this controlled run. Increasing the weight further continued to reduce PIT autocorrelation while generally worsening the other metrics.

The ten-seed financial ablation has a different ordering. There, CA-RNN improved MSE, Energy Score, interval score, nominal-coverage error, and projected-PIT calibration relative to temporal-penalty removal for every matched seed. It widened intervals, but simultaneous coverage and interval-score improvements show that the result is not merely unpenalized widening. Improvements in PIT temporal dependence were smaller and generally unresolved. Taken together, the experiments show that the penalty changes serial calibration behavior and can repair costs induced by marginal calibration in financial data, but does not provide a general predictive improvement or establish PIT independence.

5.1.5 Misspecification and distribution shift

The replicated scale-only experiment completed all 50 independent DGP realizations and 100 neural fits without recorded failures. Table A.4 reports paired CA-RNN-minus-RNN differences, with Monte Carlo standard errors in parentheses. At zero shift, the Energy Score difference is essentially zero: -0.00011 with MCSE 0.00079 . Its interval-score difference is $+0.00311$ (MCSE 0.00435), and analytic projected-PIT calibration differs by -0.000025 (MCSE 0.000024). This is nominal near-parity, not evidence of uniform CA-RNN dominance: its zero-shift test NLL is higher by 0.00629 (MCSE 0.00277) and its variogram score is higher by 0.000414 (MCSE 0.000107).

At severity 0.5, CA-RNN has lower Energy Score by 0.01624 (MCSE 0.00152), interval score by 0.29711 (MCSE 0.01749), variogram score by 0.00764 (MCSE 0.00040), analytic projected-PIT calibration error by 0.00124 (MCSE 0.000065), and test NLL by 0.22506 (MCSE 0.01534). At severity 1, the corresponding advantages increase to

0.02924, 0.55844, 0.01543, 0.00178, and 0.57664, respectively. Component-averaged MSE also favors CA-RNN under both shifts, but by smaller amounts: -0.00644 (MCSE 0.00260) and -0.01485 (MCSE 0.00429). All absolute proper scores and calibration errors worsen as scale rises for both methods; the finding is less degradation for CA-RNN, not immunity to shift. The signs of the principal shifted effects persist in the 10-, 25-, and 50-realization prefixes, although the earliest prefix gives somewhat larger effect magnitudes. No inference is made about correlation-only, tail, or mean-dynamics shifts from this experiment.

The frozen-model experiment jointly increased innovation scale and common correlation from severity zero to one. All specifications deteriorated. At zero, RNN had Energy Score 1.2022, interval score 3.8595, and projected calibration error 0.000866, compared with 1.2039, 3.9311, and 0.001345 for CA-RNN. At maximum severity, the ordering reversed: RNN obtained 2.3789, 10.5168, and 0.009528, while CA-RNN obtained 2.3474, 10.0759, and 0.007717. This single-seed combined stress is exploratory evidence rather than an isolated causal test of scale or dependence. Absolute levels and paired degradation are considered together because a worse nominal starting point can mechanically reduce measured deterioration.

5.1.6 Financial results

Table [A.5](#) reports averages over ten matched optimization seeds and 957 forecast origins on the original financial test period. CA-RNN reduced projected-PIT calibration error relative to RNN for all ten seeds, from 0.001438 to 0.001225, or approximately 15%. Its Energy Score and RMSE were worse for all ten seeds. Its mean interval score was

slightly lower, but the paired difference remained unresolved, with a 95% hierarchical interval of $[-0.0191, 0.0136]$. In contrast, its coverage-error difference was -0.00528 with interval $[-0.00815, -0.00136]$, favoring CA-RNN.

The temporal-penalty-removal ablation isolates the contribution of that objective component. Against RNN it had worse MSE, Energy Score, interval score, and coverage error despite lower projected-PIT error. CA-RNN then reduced MSE by 0.00465, Energy Score by 0.00317, interval score by 0.03749, and coverage error by 0.01032 relative to the temporal-penalty-removal ablation; every corresponding hierarchical interval excluded zero. It repairs a substantial part, but not all, of the cost of marginal calibration.

5.1.7 Post-protocol financial evaluation audit

The repeated original-panel fits were reevaluated with the exact projected Gaussian PIT and held-out joint NLL. Table A.6 separates these audit outcomes from the frozen sample-rank diagnostics. Analytic projected calibration improves from 0.001439 for RNN to 0.001223 for CA-RNN, with the latter lower for all ten optimization seeds. This agrees closely with the sample-rank ordering (0.001438 versus 0.001225), so the finite forecast-draw PIT approximation does not explain the developmental calibration effect. Nevertheless, CA-RNN has higher mean test NLL (5.3095 versus 5.2713), component-averaged MSE (0.8986 versus 0.8833), and Energy Score (1.1531 versus 1.1450). The variogram scores are nearly tied (0.156705 versus 0.156821), giving no clear dependence-sensitive proper-score advantage to either method on the original panel. Analytic projected-PIT mean absolute autocorrelation is higher for CA-RNN

(0.03691 versus 0.03072), despite the temporal penalty. The audit thus strengthens, rather than reverses, the calibration–predictive-quality distinction.

Using saved matched forecasts, a post-protocol block-length check repeated the hierarchical seed-and-origin bootstrap with blocks of 10, 20, and 40 trading days. The sign and zero-exclusion decision were stable across these lengths for MSE (CA–RNN worse), Energy Score (worse), and sample-rank projected-PIT calibration error (better). Interval-score and variogram-score differences remained unresolved at all three lengths. Temporal PIT-dependence inference was less reliable: the zero-exclusion decision changed between 10-day and longer blocks, and the block-resampled interval was displaced from the observed point difference. Artificial block boundaries can perturb lagged covariance, so this bootstrap is not used as decisive evidence for temporal PIT adequacy. These sensitivity analyses do not alter the frozen external-panel estimands or results.

5.1.8 Conventional-baseline results

Table A.7 adds the strongest traditional benchmark to the original panel. Validation selected VAR lag one, and all four marginal GARCH optimizations converged. The fixed-covariance Gaussian VAR was poorly calibrated; adding marginal GARCH scales and resampling joint standardized residuals reduced its interval score from 4.3993 to 4.0298 and projected-PIT error from 0.003388 to 0.000453. VAR–GARCH had better interval score, coverage error, and projected-PIT calibration than CA–RNN. It also had slightly better RMSE and Energy Score. The comparison attributes much of the neural models’ remaining uncertainty error to unmodeled conditional volatility rather

than an insufficiently flexible conditional mean alone.

Against CA-RNN, the paired hierarchical intervals favored VAR–GARCH for MSE, Energy Score, interval score, coverage error, and projected-PIT calibration. These intervals repeat the fixed traditional forecast across neural seed slots; they represent test-origin dependence and neural optimization variation but condition on the fitted baseline.

The training-sample specification checks in Table A.8 support the numerical and volatility adequacy of the benchmark. The VAR companion spectral radius was 0.1513, well below one. Every GARCH optimization converged, all persistence estimates were below one, and no filtered variance reached the numerical floor. Persistence was nevertheless high, ranging from 0.9700 to 0.9894, so volatility shocks decay slowly. None of the 12 Ljung–Box checks on squared standardized residuals rejected at the 5% level, providing no evidence of remaining volatility dependence at the tested lags.

The conditional-mean diagnostics are less favorable. Six of the 12 Ljung–Box tests on standardized residuals rejected at 5%; XOM rejected at all three tested lags, WMT at lags 10 and 20, and JPM at lag 5. Thus VAR(1) leaves some linear serial dependence even though validation preferred it to the longer-lag candidates. All four Jarque–Bera tests strongly rejected Gaussian standardized innovations, and marginal kurtosis ranged from 4.88 to 11.54. Pairwise standardized-residual correlations ranged from 0.22 to 0.49. The heavy tails and cross-sectional dependence jointly support the use of synchronized empirical residual vectors rather than Gaussian or independently resampled innovations. Accordingly, VAR–GARCH is treated as a strong volatility-

aware benchmark, not a fully specified model of the conditional mean.

5.1.9 Frozen external-panel replication

The SHA-256 hash recorded with the executed results matches the frozen protocol. The replication contains ten paired optimization seeds and 989 test origins; all marginal GARCH optimizations converged. Table [A.9](#) reports model means.

The frozen table retains the protocol’s sample-rank projected PIT diagnostic so that the prespecified primary outcome is not replaced after inspection. Exact Gaussian projected PIT and test-NLL audits are reported for the development-panel RNN and CA-RNN and for the replicated isolated-shift experiment, where the conditional Gaussian parameters were retained. They are not reconstructed from the external panel’s saved forecast draws: samples alone do not identify the exact predictive density. Likewise, no NLL is assigned to the empirical VAR-GARCH residual bootstrap because it does not supply the same tractable parametric density. This is a comparability boundary, not a zero or missing performance result.

For the prespecified primary contrast, CA-RNN reduced projected-PIT calibration error by 0.0000473, with 95% interval $[-0.000163, -0.000012]$. This targeted effect therefore replicated on disjoint assets. It did not yield a better predictive distribution overall: candidate-minus-RNN differences were +0.01476 for MSE, +0.01007 for Energy Score, and +0.02318 for interval score, with respective intervals $[0.01236, 0.01751]$, $[0.00847, 0.01189]$, and $[0.00773, 0.03937]$. The coverage-error difference was unresolved. Projected PIT mean absolute autocorrelation was also worse by 0.00370, with interval $[0.00005, 0.00455]$. Thus the sequential objective did not improve out-of-sample

temporal PIT adequacy despite targeting finite-lag moments during training.

The secondary comparison again favored VAR–GARCH over CA-RNN for MSE, Energy Score, interval score, coverage error, and projected calibration; all five intervals excluded zero. The RNN retained better MSE and Energy Score than VAR–GARCH, whereas VAR–GARCH had better interval score, coverage error, and projected calibration. Conditional-mean accuracy and uncertainty calibration are therefore supplied by different model components in this application.

5.1.10 Rolling market-regime results

Table A.10 gives five-seed means. CA-RNN had lower projected-PIT calibration error than RNN in all four regimes. Hierarchical intervals favored it in the pre-pandemic, inflation/tightening, and recent windows; the pandemic difference was unresolved. In contrast, its Energy Score was lower only pre-pandemic, where the difference was unresolved.

The interval-score ordering is regime-dependent. CA-RNN significantly outperformed RNN during the pandemic and inflation/tightening periods, was unresolved pre-pandemic, and significantly underperformed recently. Coverage error similarly favored CA-RNN pre-pandemic and during inflation/tightening, was unresolved during the pandemic, and favored RNN recently. These reversals directly answer the distribution-shift question.

Across regimes, CA-RNN improved Energy Score and interval score over the temporal-penalty-removal ablation directionally in four of four windows. These improvements

were resolved in the pandemic, inflation/tightening, and recent windows, while the pre-pandemic differences were unresolved. Projected calibration improved over the temporal-penalty-removal ablation in three of four windows.

5.1.11 Interval-score mechanism under temporal shift

To explain the regime reversal, the retained forecasts were analyzed without refitting. For a 90% interval $[L_{t,j}, U_{t,j}]$, the componentwise interval score decomposes exactly as

$$S_{t,j} = \underbrace{U_{t,j} - L_{t,j}}_{\text{width}} + \underbrace{20(L_{t,j} - Y_{t,j})\mathbf{1}\{Y_{t,j} < L_{t,j}\}}_{\text{lower-tail penalty}} + \underbrace{20(Y_{t,j} - U_{t,j})\mathbf{1}\{Y_{t,j} > U_{t,j}\}}_{\text{upper-tail penalty}}. \quad (5.1)$$

Table A.11 reports the mean CA-RNN-minus-RNN difference in each component. The arithmetic identity in (5.1) means the final column is exactly the sum of the preceding three columns.

The pandemic and inflation/tightening improvements arise through beneficial widening: the added width is more than offset by reductions in both tail-miss penalties. In the recent period, CA-RNN instead produces slightly narrower intervals and larger penalties in both tails. Before the pandemic, the mechanism is different again: reduced width outweighs the additional tail penalties, yielding a small and statistically unresolved sharpness benefit.

The stress-period effect is concentrated in forecasts preceded by elevated volatility. CA-RNN's interval-score differences in the pandemic are +0.0198, −0.0725, and −0.1705 for low-, medium-, and high-volatility origins, respectively. During inflation/tightening they are +0.0209, −0.0562, and −0.1045. By contrast, the recent-period differences are

positive in every state: +0.0318, +0.0238, and +0.0305. Volatility states use trailing 20-day realized volatility known before each forecast and within-regime terciles.

The improvements are also asset-concentrated. XOM accounts for most of the inflation/tightening gain, with an interval-score difference of -0.2364 ; AAPL, JPM, and WMT have differences of -0.0051 , $+0.0221$, and $+0.0331$. Pandemic improvements occur for AAPL, XOM, and WMT but not JPM. Recent losses are concentrated in AAPL, XOM, and WMT, while JPM improves slightly. These results preclude interpreting an aggregate gain as uniform across assets.

The train-to-test descriptors support, but do not prove, a scale-and-tail interpretation. Pandemic volatility is 1.26 times its training level and the mean standardized fourth moment is 29.87. The corresponding values are 1.00 and 12.08 during inflation/tightening, versus 0.83 and 6.37 recently. The recent period has the largest standardized location shift and correlation shift of the four windows. Thus the complete CA-RNN objective appears most useful for persistent scale and tail errors, while lower-volatility location or dependence changes remain harder.

Importantly, projected and coordinate PIT calibration still improve recently even as 90% interval performance deteriorates. The calibration loss averages 19 CDF grid levels and multiple projections; it does not target the 5% and 95% quantiles specifically. Better global PIT uniformity can therefore coexist with worse tail coverage. The exact score decomposition is established by the retained forecasts, whereas the proposed economic mechanism is post-hoc and requires new data for confirmation.

5.2 Discussion

5.2.1 Interpretation

Four conclusions follow. First, a marginal calibration penalty is insufficient for dependent multivariate forecasting: it can improve its targeted diagnostic while harming proper scores and nominal coverage. Second, the temporal PIT-moment loss produces a genuine but narrow intervention on serial PIT moments. It recovered part of the temporal-penalty-removal model’s predictive-performance loss in the ten-seed financial experiment, but the refreshed controlled simulation found only a small temporal-dependence reduction accompanied by worse point error, proper scores, projected calibration, and coverage error. Its value is therefore data- and regime-dependent rather than automatic. Third, CA-RNN’s projected-calibration effect transfers to a disjoint asset panel, but its developmental interval and coverage improvements do not: the frozen replication instead favors RNN on all three proper or accuracy-oriented scores and finds no resolved coverage benefit. Fourth, VAR-GARCH shows that explicit volatility dynamics and joint residual resampling can provide better-calibrated uncertainty than direct calibration regularization in these data.

The replicated isolated-scale experiment provides a complementary and more specific robustness finding. Across 50 independent Gaussian DGP realizations, CA-RNN offers little nominal proper-score advantage but degrades less than RNN when only test innovation scale rises. Its shifted gains span joint NLL, Energy Score, variogram score, marginal interval score, and analytic projected-PIT calibration, rather than a single

diagnostic. This does not contradict the frozen financial replication: the simulation isolates one known perturbation of a Gaussian autoregression, while the external asset panel combines unknown volatility, tail, dependence, and mean changes. The post-protocol simulation therefore establishes a boundary at which the fixed CA-RNN objective helps, not a general claim of financial forecast superiority.

The disagreement between Energy Score and interval score is substantively useful. Energy Score assesses the multivariate predictive distribution, whereas the marginal interval score rewards coverage and sharpness asset by asset. During the pandemic and inflation/tightening periods, CA-RNN improves the latter without improving the former. Its practical value therefore depends on whether an application prioritizes marginal risk intervals or overall multivariate accuracy. The confirmatory evidence narrows this claim further: lower projected-PIT grid error alone is insufficient evidence of improved operational uncertainty.

The regret identities explain the observed divergence. Energy Score regret is one half of the energy distance and therefore responds to discrepancies anywhere in the joint distribution. Marginal interval-score regret is the sum of two integrated quantile errors at the selected tail probabilities. By contrast, the training regularizer averages a finite set of smoothed PIT-CDF moments across grid points and projections. Because this finite criterion is non-identifying, reducing it supplies no mathematical upper bound on either regret. The external result—lower projected-PIT error alongside higher Energy and interval scores—is therefore compatible with the objectives rather than paradoxical.

5.2.2 Limitations

The central limitation concerns what the calibration objective can identify. CA-RNN evaluates a finite set of smoothed projected-PIT moments, estimated from dependent mini-batches. Its conclusions therefore depend on the selected projections, grid, smoothing parameter, and batching scheme. Even a small value of this empirical loss cannot establish full joint or history-conditional calibration. The temporal term narrows one part of this gap by targeting selected lagged covariances, but finite zero autocovariances do not imply PIT independence. Moving-block resampling introduces an additional complication because artificial block boundaries can affect autocorrelation estimates. The 10/20/40-day sensitivity audit changed the zero-exclusion decision for PIT dependence and showed resampling bias relative to its observed point difference; the corresponding uncertainty intervals are therefore interpreted as diagnostic rather than definitive.

The empirical scope is also deliberately limited. The financial evidence comes from two four-asset panels and one-step-ahead forecasts, while the chronological regimes are calendar-based and were partly formulated after broader exposure to the development data. The frozen external panel reduces, but does not eliminate, dependence on choices of assets, architecture, tuning path, projections, and evaluation windows. Similarly, the hierarchical bootstrap represents variation across neural optimization seeds and dependent forecast origins conditional on those design choices. Because each conventional benchmark is fitted once per panel, its reported comparisons do not integrate over baseline parameter- estimation uncertainty. The isolated-scale simulation adds 50 independent datasets, but uses one neural seed per dataset and was

designed after earlier developmental findings. Its MCSE therefore does not separate DGP from optimization variation or protect against post-selection of the shift family.

Finally, the neural predictive law is Gaussian and does not explicitly encode the persistent conditional volatility or heavy tails characteristic of financial returns. This matters for interpretation: VAR–GARCH’s advantage may reflect a better uncertainty specification rather than a general advantage of linear over recurrent conditional-mean dynamics. The original-panel benchmark is itself not fully dynamically specified. Its diagnostics are in-sample and descriptive, several Ljung–Box tests on standardized residual levels reject, and the p -values are not adjusted for fitted parameters or multiple testing. Its strong performance is therefore attributed primarily to effective volatility filtering and joint empirical innovation resampling, not to complete model adequacy.

5.2.3 Future work

The most direct methodological extension is a tail-aware sequential calibration objective. The current loss weights all 19 PIT grid levels and all fixed projections equally. A weighted version could emphasize application-relevant tail levels,

$$\mathcal{L}_{\text{tail}} = \frac{1}{R \sum_g w_g} \sum_{r=1}^R \sum_{g=1}^G w_g \left[\frac{1}{|B|} \sum_{t \in B} s_\tau(q_g - U_{t,r}) - q_g \right]^2, \quad (5.2)$$

with larger prespecified w_g near $q = 0.05$ and $q = 0.95$. A complementary extension would penalize lagged dependence in tail exceedance indicators rather than only covariance of centered PITs. Either extension must be tuned on validation data and evaluated with global calibration and proper scores to avoid improving one nominal interval at the expense of the rest of the distribution.

A particularly relevant extension is a volatility-aware CA-RNN. The recurrent state could parameterize time-varying marginal scales and a structured dynamic correlation matrix, with Student- t or mixture innovations replacing the current Gaussian head. This would test whether calibration-aware estimation adds value after conditional heteroskedasticity and heavy tails have been modeled directly, rather than requiring the calibration penalty to compensate for their omission. It is intentionally left outside the present study because a defensible evaluation would require a factorial comparison of volatility-aware RNN and volatility-aware CA-RNN, alongside their current Gaussian counterparts. Without that control, improvements could not be attributed separately to volatility modeling, the innovation distribution, or calibration-aware training.

Further work should also test new frequencies, forecast horizons, broader asset universes, and genuinely future calendar periods. Conditional calibration diagnostics could use volatility, market return, and liquidity states fixed before evaluation. Finally, uncertainty for PIT dependence should use a procedure that avoids correlations introduced at moving-block boundaries.

5.3 Conclusion

This paper adapts calibration-aware Bayesian learning to continuous multivariate recurrent forecasting and introduces a temporal calibration penalty based on lagged projected-PIT covariance. Calibration-aware training reliably changes forecast calibration, but projected calibration alone incurs predictive costs. The temporal PIT-moment component improves the corresponding loss-component ablation in the ten-seed financial study, but not across every metric in the controlled simulation. The complete

CA-RNN nevertheless produces lower projected-PIT calibration error than a likelihood-trained RNN across financial optimization seeds, chronological regimes, and a frozen disjoint-asset replication; the same ordering does not hold in the single-realization nonlinear simulation.

The method is not uniformly superior. In the confirmatory experiment it has worse MSE, Energy Score, interval score, and temporal PIT dependence than RNN, while its coverage-error difference is unresolved. VAR-GARCH also outperforms it on the principal uncertainty metrics in both asset panels. The contribution is therefore not a dominance claim, nor evidence that the temporal PIT-moment penalty achieves PIT independence. It is a controlled demonstration that calibration-aware training can reproducibly improve its targeted finite-projection diagnostic while failing to improve—and sometimes worsening—proper scores and sequential adequacy. In the completed post-protocol isolated-scale experiment, 50 paired realizations show a more precise robustness crossover: nominal near-parity gives way to lower NLL, Energy Score, variogram score, interval score, and analytic projected-PIT error for CA-RNN as test innovation scale rises. This is evidence for one isolated stress mechanism, not general dominance under distribution shift. Strong volatility-aware baselines, ablation across data-generating settings, and external confirmation are thus essential when evaluating calibration-aware neural forecasts.

Bibliographic references

1. Huang, J., Park, S. & Simeone, O. Calibration-aware bayesian learning. *arXiv preprint arXiv:2305.07504* (2024).
2. Utpala, S. & Rai, P. Quantile regularization: towards implicit calibration of regression models. *arXiv preprint arXiv:2002.12860* (2020).
3. Dheur, J. & Ben Taieb, S. *A large-scale study of probabilistic calibration in neural network regression in Proceedings of the 40th International Conference on Machine Learning* **202** (2023), 7813–7836.
4. Laajil, A., Zhalieva, E., Desobry, N. & Ben Taieb, S. Calibrated multivariate distributional regression with pre-rank regularization. *arXiv preprint arXiv:2601.22895* (2026).
5. Gneiting, T., Balabdaoui, F. & Raftery, A. E. Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society: Series B* **69**, 243–268. doi:[10.1111/j.1467-9868.2007.00587.x](https://doi.org/10.1111/j.1467-9868.2007.00587.x) (2007).
6. Gneiting, T. & Raftery, A. E. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association* **102**, 359–378. doi:[10.1198/016214506000001437](https://doi.org/10.1198/016214506000001437) (2007).
7. Gneiting, T., Stanberry, L. I., Grimit, E. P., Held, L. & Johnson, N. A. Assessing probabilistic forecasts of multivariate quantities, with an application to ensemble predictions of surface winds. *TEST* **17**, 211–235. doi:[10.1007/s11749-008-0114-x](https://doi.org/10.1007/s11749-008-0114-x) (2008).
8. Scheuerer, M. & Hamill, T. M. Variogram-based proper scoring rules for probabilistic forecasts of multivariate quantities. *Monthly Weather Review* **143**, 1321–1334. doi:[10.1175/MWR-D-14-00269.1](https://doi.org/10.1175/MWR-D-14-00269.1) (2015).
9. Gneiting, T. & Resin, J. Regression diagnostics meets forecast evaluation: conditional calibration, reliability diagrams, and coefficient of determination. *Electronic Journal of Statistics* **17**, 3226–3286. doi:[10.1214/23-EJS2180](https://doi.org/10.1214/23-EJS2180) (2023).

Bibliographic references

10. Diebold, F. X., Gunther, T. A. & Tay, A. S. Evaluating density forecasts, with applications to financial risk management. *International Economic Review* **39**, 863–883 (1998).
11. Berkowitz, J. Testing density forecasts, with applications to risk management. *Journal of Business & Economic Statistics* **19**, 465–474. doi:[10.1198/07350010152596718](https://doi.org/10.1198/07350010152596718) (2001).
12. Kn"uppel, M. Evaluating the calibration of multi-step-ahead density forecasts using raw moments. *Journal of Business & Economic Statistics* **33**, 270–281 (2015).
13. Guo, C., Pleiss, G., Sun, Y. & Weinberger, K. Q. *On calibration of modern neural networks* in *Proceedings of the 34th International Conference on Machine Learning* **70** (2017), 1321–1330.
14. Kumar, A., Sarawagi, S. & Jain, U. *Trainable calibration measures for neural networks from kernel mean embeddings* in *Proceedings of the 35th International Conference on Machine Learning* **80** (2018), 2805–2814.
15. Kuleshov, V., Fenner, N. & Ermon, S. *Accurate uncertainties for deep learning using calibrated regression* in *Proceedings of the 35th International Conference on Machine Learning* **80** (2018), 2796–2804.
16. Dheur, J. & Ben Taieb, S. *Probabilistic calibration by design for neural network regression* in *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics* **238** (2024), 3133–3141.
17. Chung, Y., Neiswanger, W., Char, I. & Schneider, J. *Beyond pinball loss: quantile methods for calibrated uncertainty quantification* in *Advances in Neural Information Processing Systems* **34** (2021).
18. Allen, S., Ziegel, J. F. & Ginsbourger, D. Assessing the calibration of multivariate probabilistic forecasts. *Quarterly Journal of the Royal Meteorological Society* **150**, 1315–1335. doi:[10.1002/qj.4647](https://doi.org/10.1002/qj.4647) (2024).
19. Desobry, N., Zhalieva, E. & Ben Taieb, S. Enforcing calibration in multi-output probabilistic regression with pre-rank regularization. *arXiv preprint arXiv:2510.21273* (2025).

Bibliographic references

20. Salinas, D., Flunkert, V., Gasthaus, J. & Januschowski, T. DeepAR: probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting* **36**, 1181–1191. doi:[10.1016/j.ijforecast.2019.07.001](https://doi.org/10.1016/j.ijforecast.2019.07.001) (2020).
21. White, H. Maximum likelihood estimation of misspecified models. *Econometrica* **50**, 1–25. doi:[10.2307/1912526](https://doi.org/10.2307/1912526) (1982).
22. Ovadia, Y. *et al.* Can you trust your model’s uncertainty? evaluating predictive uncertainty under dataset shift in *Advances in Neural Information Processing Systems* **32** (2019).
23. Bollerslev, T. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics* **31**, 307–327. doi:[10.1016/0304-4076\(86\)90063-1](https://doi.org/10.1016/0304-4076(86)90063-1) (1986).

Appendix A

Tables and Figures

A.1 Tables

All tables cited in the main text are collected here in order of first citation.

Table A.1: Positioning relative to the closest calibration literature. “Shift” means an explicit out-of-distribution or regime evaluation, not ordinary held-out testing.

Work	Outcome	Training penalty	Multivariate pre-rank	Sequential/shift focus
Huang et al. [1]	Classes	Differentiable ECE	No	No
Dheur–Ben Taieb [3]	Scalar regression	PIT-CDF	No	No
Laajil et al. [4]	Vector regression	PIT-CDF	Yes	No; i.i.d. design
This study	Vector time series	PIT-CDF plus lag moments	Fixed linear pre-ranks	Yes

Table A.2: Mapping from calibration-aware classification to forecasting.

Source-paper model	Forecasting model	Objective components
FNN	RNN	NLL
BNN	BRNN	expected NLL + KL
CA-FNN	CA-RNN	NLL + projected-PIT calibration
CA-BNN	CA-BRNN	expected NLL/calibration + KL

Table A.3: Loss-component ablation.

Specification	NLL	$\mathcal{L}_{\text{cal}}$	$\mathcal{L}_{\text{seq}}$
RNN	✓		
CA-RNN (temporal penalty removed)	✓	✓	
CA-RNN	✓	✓	✓

Table A.4: Isolated Gaussian innovation-scale shift over 50 independent paired realizations. Each entry is the mean CA-RNN-minus-RNN difference with its realization-level Monte Carlo standard error in parentheses; negative values favor CA-RNN.

Severity	Energy Score	Variogram	Interval score	Analytic PIT error
0	−0.00011 (0.00079)	+0.000414 (0.000107)	+0.00311 (0.00435)	−0.000025 (0.000024)
0.5	−0.01624 (0.00152)	−0.007638 (0.000400)	−0.29711 (0.01749)	−0.001241 (0.000065)
1	−0.02924 (0.00215)	−0.015428 (0.000766)	−0.55844 (0.02850)	−0.001781 (0.000090)

Table A.5: Mean financial performance over ten matched optimization seeds. Lower is better except for coverage, whose nominal target is 0.90.

Model	RMSE	Energy	Interval	Coverage	PIT calibration
RNN	0.9399	1.1450	4.0855	0.8717	0.001438
CA-RNN (temporal penalty removed)	0.9504	1.1563	4.1191	0.8666	0.001306
CA-RNN	0.9479	1.1531	4.0816	0.8770	0.001225

Table A.6: Original-panel post-protocol evaluation audit: means across ten matched neural seeds. All entries are losses or diagnostic discrepancies for which lower is better. MSE and RMSE are distinct summaries.

Model	MSE	RMSE	Test NLL	Variogram	Analytic PIT error
RNN	0.8833	0.9399	5.2713	0.156821	0.001439
CA-RNN	0.8986	0.9479	5.3095	0.156705	0.001223

Table A.7: Original-panel neural and VAR–GARCH results. Neural entries are ten-seed means; the traditional model is fitted once. Coverage targets 0.90.

Model	RMSE	Energy	Interval	Coverage	PIT calibration
RNN	0.9399	1.1450	4.0855	0.8717	0.001438
CA-RNN	0.9479	1.1531	4.0816	0.8770	0.001225
VAR–GARCH bootstrap	0.9452	1.1495	4.0298	0.8989	0.000453

Table A.8: Training-sample diagnostics for the selected VAR(1)–GARCH benchmark. $p_{z,20}$ and $p_{z^2,20}$ are descriptive Ljung–Box p -values at lag 20; kurtosis is the ordinary, rather than excess, standardized-residual kurtosis.

Asset	$\alpha + \beta$	$p_{z,20}$	$p_{z^2,20}$	Kurtosis
AAPL	0.9700	0.2307	0.9689	6.30
JPM	0.9849	0.1490	0.7370	5.53
XOM	0.9849	0.0142	0.8253	4.88
WMT	0.9894	0.0089	0.8922	11.54

Table A.9: Frozen MSFT–BAC–CVX–COST replication. Neural entries are ten-seed means; coverage targets 0.90.

Model	RMSE	Energy	Interval	Coverage	PIT cal.	PIT dep.
RNN	0.8609	1.0677	3.7766	0.8703	0.000783	0.03008
CA-RNN	0.8694	1.0778	3.7998	0.8683	0.000735	0.03378
VAR–GARCH bootstrap	0.8646	1.0713	3.7528	0.8928	0.000491	0.03270

Table A.10: Five-seed regime means for RNN (R) and CA-RNN (C).

Regime	Energy Score		Interval score		PIT calibration	
	R	C	R	C	R	C
Pre-pandemic	0.8629	0.8626	3.1123	3.1027	0.001914	0.001503
Pandemic	1.5016	1.5029	5.7371	5.6629	0.002995	0.002883
Inflation/tightening	1.2809	1.2884	4.5862	4.5396	0.001842	0.001465
Recent	1.0237	1.0289	3.6648	3.6935	0.001224	0.000904

Table A.11: Decomposition of the interval-score difference between CA-RNN and RNN. Negative total differences favor CA-RNN.

Regime	Width	Lower penalty	Upper penalty	Total
Pre-pandemic	−0.0439	+0.0241	+0.0103	−0.0096
Pandemic	+0.0420	−0.0599	−0.0563	−0.0742
Inflation/tightening	+0.2342	−0.1709	−0.1098	−0.0466
Recent	−0.0185	+0.0268	+0.0205	+0.0287

A.2 Figures

All figures cited in the main text are collected here in order of first citation.




Figure A.1: Paired CA–RNN-minus-RNN differences across scale severity. Error bars show ± 1.96 Monte Carlo standard errors across independent DGP realizations; they do not cover uncertainty over alternative DGP families or tuning choices.

Appendix A. Tables and Figures




Figure A.2: Mean Energy Score, interval score, and projected-PIT calibration error across five matched seeds and four chronological test regimes. Lower is better.




Figure A.3: Exact decomposition of CA-RNN-minus-RNN interval-score differences. Tail reductions offset widening in the pandemic and inflation/tightening periods, but not in the recent period.

Appendix B

Reproducibility and Post-Protocol Evaluation Audit

B.1 Data, transformations, and chronological partitions

The development panel contains adjusted daily closing prices for AAPL, JPM, XOM, and WMT, requested for 2007-01-01 through 2026-01-01 (exclusive). After complete-case date alignment and one-day log differencing, it contains 4,779 return observations from 2007-01-04 through 2025-12-31. The 60/20/20 chronological split gives training (2,867 observations, 2007-01-04 to 2018-05-23), validation (955, 2018-05-24 to 2022-03-09), and test (957, 2022-03-10 to 2025-12-31). The observed adjusted-price span begins 2007-01-03. The cleaning rule rejects any ticker with over 1% missingness, then retains complete dates and positive prices. The panel is subject to survivorship-biased asset selection.

The frozen external panel contains MSFT, BAC, CVX, and COST in that order. Its 4,945 aligned log-return observations run from 2007-01-04 through 2026-08-31. The same split gives training (2,967 observations, 2007-01-04 to 2018-10-15), validation (989, 2018-10-16 to 2022-09-20), and test (989, 2022-09-21 to 2026-08-31). Each component is standardized using its own training-partition mean and sample standard deviation. A 20-observation historical context precedes each validation and test target; neither scaler nor fitted model sees the target or later observations during selection.

B.2 Model, randomization, and selection settings

The focused neural architecture is a one-layer, 32-hidden-unit GRU with a full-covariance Gaussian predictive head and minimum Cholesky diagonal scale 10^{-4} . It uses a 20-day input history, AdamW learning rate 10^{-3} , batch size 64, maximum 200 epochs, patience 25, and validation-NLL checkpoint selection in the repeated financial experiments. The CA-RNN loss weights are $\lambda_{\text{cal}} = 100$ and $\lambda_{\text{seq}} = 2000$; the PIT CDF grid has 19 points, smoothing temperature 0.03, and five temporal lags. Projection directions comprise four coordinates, an equal-weight direction, and four random unit vectors generated with seed 606. The repeated financial and frozen external studies use neural training seeds 611, 619, 631, 641, 647, 653, 661, 673, 683, and 691, forecast seed 20,606, and 500 draws per forecast origin. The original-panel hierarchical bootstrap uses 20-trading-day blocks, 2,000 draws, and seed 808; the frozen panel uses the same block length and draw count with seed 1208. All paired comparisons reuse forecast-origin indices across models. The conventional benchmark selects VAR lag from $\{1, 5, 10, 20\}$ by validation one-step MSE and then fits marginal GARCH(1, 1) and the joint standardized-residual resampling law on the training partition.

The four rolling test windows are 2017–2019, 2020–2021, 2022–2023, and 2024–2025. Their expanding training endpoints are, respectively, the ends of 2014, 2016, 2019, and 2021; intervening validation periods are 2015–2016, 2017–2019, 2020–2021, and 2022–2023. The regime training seeds are 701, 709, 719, 727, and 733. Their 1,000-draw hierarchical bootstraps are developmental rather than frozen confirmatory analyses. The four-model nonlinear demonstration instead uses 1,200 observations, DGP and training seed 2027, 150 maximum epochs, and patience 20. The original mixed scale-

and-correlation curve uses DGP and training seed 505, 120 maximum epochs, and patience 15; its single-realization interpretation remains exploratory.

The completed post-protocol isolated-scale experiment is implemented in notebook 14 and reported in Table A.4. It uses 50 independent Gaussian autoregressive realizations with DGP seeds 1000–1049, one distinct neural seed $606+b$ per realization $b = 0, \dots, 49$, and common forecast seed 20,606. Each series has 1,200 observations; the first 60% train, the next 20% validate, and the final 20% test. The shift begins exactly at the test boundary. Severities $\{0, 0.5, 1\}$ multiply the post-boundary innovation standard deviation by $1 + \delta$ while holding correlation, autoregressive coefficients, and innovation family fixed. RNN and CA-RNN are fitted once per realization with 120 maximum epochs, patience 15, and validation-NLL checkpointing on the same unshifted training/validation history, then held fixed over severities. The run records failures and training time. For each severity, paired CA-RNN-minus-RNN differences are averaged across valid realizations, with Monte Carlo standard error $s_D/\sqrt{B}$. Prefixes of 10, 25, and 50 realizations provide a replication-count sensitivity check. One neural seed per DGP realization means DGP and optimization variability are not separately identified in this experiment.

B.3 Evaluation definitions and protocol status

For a d -asset standardized target y_t and forecast mean $\hat{\mu}_t$, the reported component-averaged mean squared error is $\text{MSE} = (Td)^{-1} \sum_{t,j} (y_{t,j} - \hat{\mu}_{t,j})^2$ and $\text{RMSE} = \sqrt{\text{MSE}}$. A seed-level mean of RMSEs need not equal the square root of the corresponding seed-level mean MSE. Neural test NLL is the mean joint Gaussian negative log density,

Appendix B. Reproducibility and Post-Protocol Evaluation Audit

evaluated at held-out targets; for Bayesian models it is the negative log of a Monte Carlo predictive mixture. For a deterministic Gaussian forecast, projected PITs are evaluated analytically as $\Phi[(a^\top y_t - a^\top \mu_t)/\sqrt{a^\top \Sigma_t a}]$ in the new evaluation audit. The previously reported sample-rank PITs remain the frozen protocol measure and are not retrospectively replaced. The variogram score is reported as the mean over origins and unordered asset pairs of $(|y_{t,i} - y_{t,j}|^{1/2} - E_F|Y_i - Y_j|^{1/2})^2$; the expectation is estimated from forecast draws. Notebook 15 checks sensitivity of developmental paired intervals to block lengths 10, 20, and 40, holding the forecast arrays fixed. These additions are secondary, post-protocol audits, not new primary outcomes.

The frozen primary protocol is preserved verbatim at `docs/confirmatory_protocol.md`; its SHA-256 digest is `20a0531461992b11d702d5b72b51c57a5e9b162e84dd5f6e6e34e5fadec570b0`. Its historical “sequential CA–RNN” label denotes the unchanged full projected-plus-temporal objective now called CA–RNN. The protocol fixes the external assets and date boundary, architecture, penalties, seeds, forecast draws, primary metrics, and 20-day hierarchical bootstrap. No new analytic PIT, NLL, variogram, or sensitivity result supersedes its originally frozen sample-PIT and proper-score comparisons.